


FP32 vs mixed precision: where did GPU work change?



Profiled kernel durations, not unprofiled forward latency. Copy / conversion is not exclusively dtype conversion.